

# Exp 9 Decision Tree with Cross Validation and GridSearchCV

## Milindi Shah J057

In [1]:

```
import os
print(os.getcwd())
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

C:\Users\shahm\Desktop\5th SEM\Machine Learning\Submissions

In [2]:

```
df = pd.read_csv('car_evaluation.csv', header = None)
```

In [3]:

```
df.head()
```

Out[3]:

	0	1	2	3	4	5	6
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

In [4]:

```
col_names = ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']
df.columns = col_names
col_names
```

Out[4]:

```
['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']
```

In [5]:

```
df.head()
```

Out[5]:

	buying	maint	doors	persons	lug_boot	safety	class
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
 #   Column        Non-Null Count  Dtype  
---  -
 0   buying        1728 non-null   object 
 1   maint         1728 non-null   object 
 2   doors         1728 non-null   object 
 3   persons       1728 non-null   object 
 4   lug_boot      1728 non-null   object 
 5   safety        1728 non-null   object 
 6   class         1728 non-null   object 
dtypes: object(7)
memory usage: 94.6+ KB
```

In [7]:

```
for i in col_names:  
    print(df[i].value_counts())
```

```
low      432  
vhigh    432  
med      432  
high     432  
Name: buying, dtype: int64  
low      432  
vhigh    432  
med      432  
high     432  
Name: maint, dtype: int64  
3        432  
2        432  
4        432  
5more    432  
Name: doors, dtype: int64  
more     576  
2        576  
4        576  
Name: persons, dtype: int64  
big      576  
med      576  
small    576  
Name: lug_boot, dtype: int64  
low      576  
med      576  
high     576  
Name: safety, dtype: int64  
unacc    1210  
acc       384  
good      69  
vgood     65  
Name: class, dtype: int64
```

In [8]:

```
df.shape
```

Out[8]:

```
(1728, 7)
```

In [9]:

```
X = df.drop(['class'],axis = 1)  
y = df['class']
```

In [10]:

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,random_state=42)
```

In [11]:

```
from sklearn.preprocessing import OrdinalEncoder
enc = OrdinalEncoder()
X_train = enc.fit_transform(X_train)
X_test = enc.transform(X_test)
```

## Gini index as criterion

In [12]:

```
from sklearn.tree import DecisionTreeClassifier
```

In [13]:

```
clf_gini = DecisionTreeClassifier(criterion='gini', max_depth=3, random_state=42)
clf_gini.fit(X_train, y_train)
```

Out[13]:

```
DecisionTreeClassifier(max_depth=3, random_state=42)
```

In [14]:

```
y_pred = clf_gini.predict(X_test)
```

In [15]:

```
from sklearn.metrics import accuracy_score
print(f'Model with gini index gives an accuracy of: {accuracy_score(y_test, y_pred)}')
```

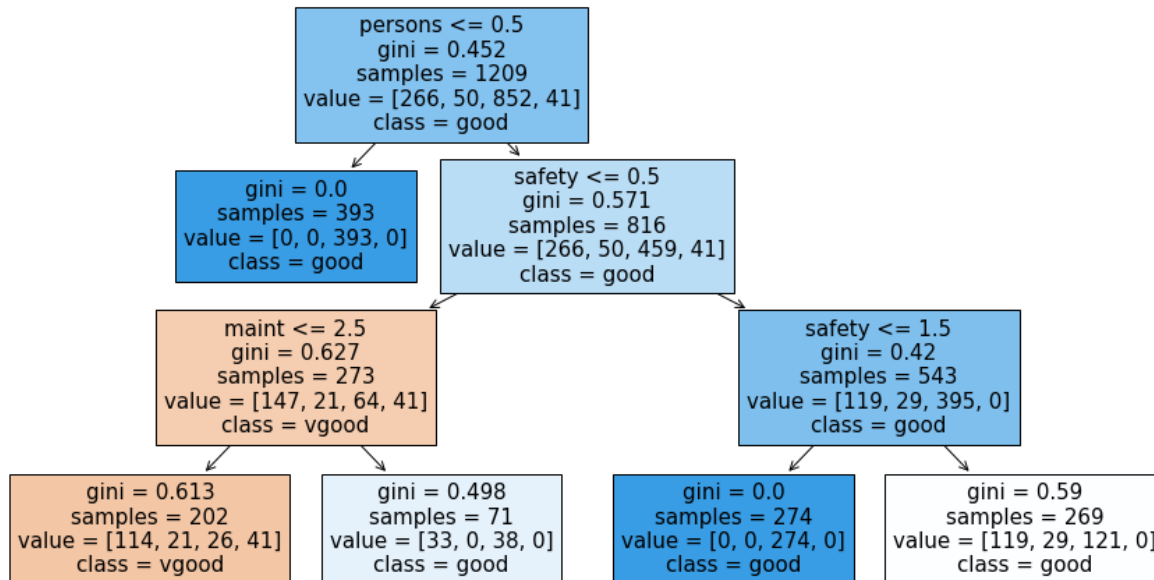
```
Model with gini index gives an accuracy of: 0.7572254335260116
```

In [16]:

```

from sklearn import tree
plt.figure(figsize=(15,8))
tree.plot_tree(clf_gini,
               feature_names=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety'],
               class_names= list(set(y_train)),
               filled = True)
plt.show()

```



In [17]:

```

# Check for underfitting

print(f'Training set score: {clf_gini.score(X_train,y_train)}')
print(f'Test set score: {clf_gini.score(X_test,y_test)}')

```

Training set score: 0.7775020678246485

Test set score: 0.7572254335260116

## Entropy as criterion

In [18]:

```

clf_entropy = DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=42)
clf_entropy.fit(X_train, y_train)

```

Out[18]:

DecisionTreeClassifier(criterion='entropy', max\_depth=3, random\_state=42)

In [19]:

```

y_pred = clf_entropy.predict(X_test)

```

In [20]:

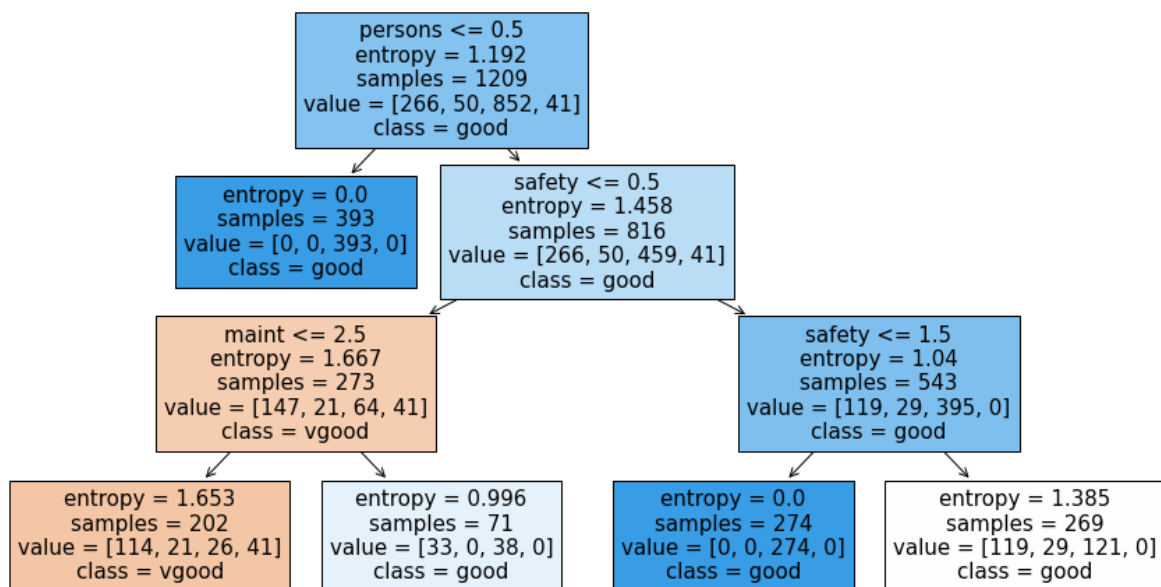
```
from sklearn.metrics import accuracy_score

print(f'Model with gini index gives an accuracy of: {accuracy_score(y_test, y_pred)}')
```

Model with gini index gives an accuracy of: 0.7572254335260116

In [21]:

```
plt.figure(figsize=(15,8))
tree.plot_tree(clf_entropy,
               feature_names=['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety'],
               class_names= list(set(y_train)),
               filled = True)
plt.show()
```



In [22]:

```
# Check for underfitting

print(f'Training set score: {clf_entropy.score(X_train,y_train)}')
print(f'Test set score: {clf_entropy.score(X_test,y_test)}')
```

Training set score: 0.7775020678246485

Test set score: 0.7572254335260116

In [23]:

```
from sklearn.metrics import confusion_matrix, classification_report
cm = confusion_matrix(y_test, y_pred)
```

In [24]:

```
print(cm)
```

```
[[ 44   0  74   0]
 [  9   0  10   0]
 [  9   0 349   0]
 [ 24   0   0   0]]
```

In [25]:

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
acc	0.51	0.37	0.43	118
good	0.00	0.00	0.00	19
unacc	0.81	0.97	0.88	358
vgood	0.00	0.00	0.00	24
accuracy			0.76	519
macro avg	0.33	0.34	0.33	519
weighted avg	0.67	0.76	0.71	519

C:\Users\shahm\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1221: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

## Cross Validation

In [36]:

```
params_grid={'criterion':['gini', 'entropy'], 'max_depth':[3,4,5,6,7,8,9,10]}
```

In [38]:

```
decision_tree=DecisionTreeClassifier()
decision_tree.fit(X_train,y_train)
```

Out[38]:

```
DecisionTreeClassifier()
```

In [48]:

```
from sklearn.model_selection import GridSearchCV
dt_validated=GridSearchCV(estimator=decision_tree,param_grid=params_grid,scoring='accuracy')
```

In [49]:

```
%%time
dt_validated.fit(X_train,y_train)
```

Wall time: 1.74 s

Out[49]:

```
GridSearchCV(cv=20, estimator=DecisionTreeClassifier(),
             param_grid={'criterion': ['gini', 'entropy'],
                          'max_depth': [3, 4, 5, 6, 7, 8, 9, 10]},
             scoring='accuracy')
```

In [53]:

```
print(f'Best parameters for decison tree classifier after CV -> {dt_validated.best_params_}')
print(f'Best score on decision tree classifier after CV -> {dt_validated.best_score_}')
```

```
Best parameters for decison tree classifier after CV -> {'criterion': 'entro
py', 'max_depth': 10}
```

```
Best score on decision tree classifier after CV -> 0.980136612021858
```

In [55]:

```
print(f'Score on train set of DT classifier before CV -> {decision_tree.score(X_train, y_tr
print(f'Score on test set of DT classifier before CV -> {decision_tree.score(X_test, y_test
print(f'Score on train set of DT classifier after CV -> {dt_validated.score(X_train, y_tra
print(f'Score on test set of DT classifier after CV -> {dt_validated.score(X_test, y_test))}
```

```
Score on train set of DT classifier before CV -> 1.0
```

```
Score on test set of DT classifier before CV -> 0.9653179190751445
```

```
Score on train set of DT classifier after CV -> 0.9925558312655087
```

```
Score on test set of DT classifier after CV -> 0.9595375722543352
```

In [56]:

```
print('Classification report on train set')
print(classification_report(y_true=y_train,y_pred=dt_validated.predict(X_train)))
```

Classification report on train set

	precision	recall	f1-score	support
acc	0.98	0.99	0.99	266
good	0.98	0.98	0.98	50
unacc	1.00	0.99	1.00	852
vgood	0.98	1.00	0.99	41
accuracy			0.99	1209
macro avg	0.98	0.99	0.99	1209
weighted avg	0.99	0.99	0.99	1209



In [57]:

```
print('Classification report on test set')
print(classification_report(y_true=y_test,y_pred=dt_validated.predict(X_test)))
```

Classification report on test set

	precision	recall	f1-score	support
acc	0.92	0.92	0.92	118
good	0.71	0.89	0.79	19
unacc	0.99	0.98	0.99	358
vgood	0.88	0.88	0.88	24
accuracy			0.96	519
macro avg	0.88	0.92	0.89	519
weighted avg	0.96	0.96	0.96	519

In [58]:

```
print('Confusion matrix on train set')
print(confusion_matrix(y_true=y_train,y_pred=dt_validated.predict(X_train)))
```

Confusion matrix on train set

```
[[263  1  2  0]
 [  0 49  0  1]
 [  5  0 847  0]
 [  0  0  0 41]]
```

In [59]:

```
print('Confusion matrix on test set')
print(confusion_matrix(y_true=y_test,y_pred=dt_validated.predict(X_test)))
```

Confusion matrix on test set

```
[[108  7  2  1]
 [  0 17  0  2]
 [  6  0 352  0]
 [  3  0  0 21]]
```

In [60]:

```
best_dt=DecisionTreeClassifier(criterion='entropy',max_depth=9)
```

In [61]:

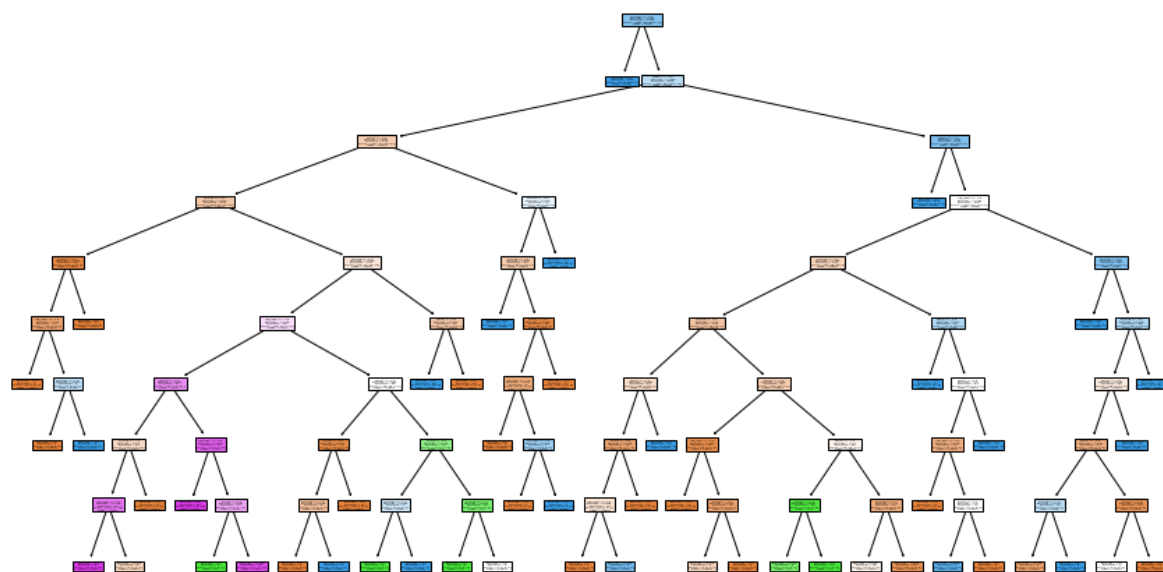
```
best_dt.fit(X_train,y_train)
```

Out[61]:

```
DecisionTreeClassifier(criterion='entropy', max_depth=9)
```

In [65]:

```
plt.figure(figsize=(15,8))  
tree.plot_tree(best_dt,feature_names=['buying','maint','doors','persons','lug_boot','safety'])  
plt.show()
```



In [ ]: